

CS202 Project

Performance evaluation and comparison between distributed file systems

Group 13

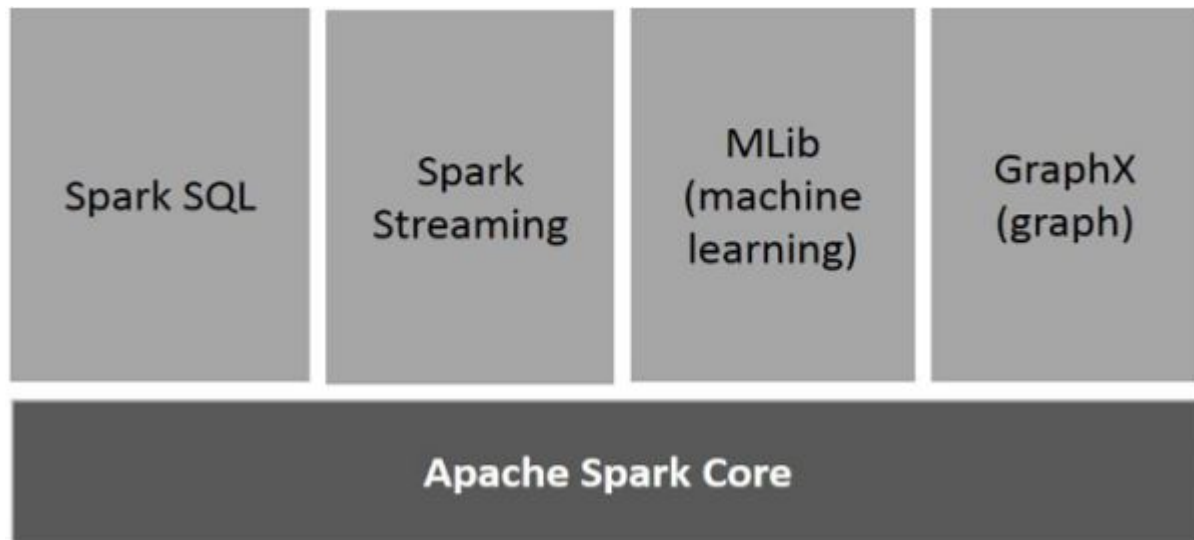
Chen-Yang Yu 862052273

Po-Cheng Kuo 862029279

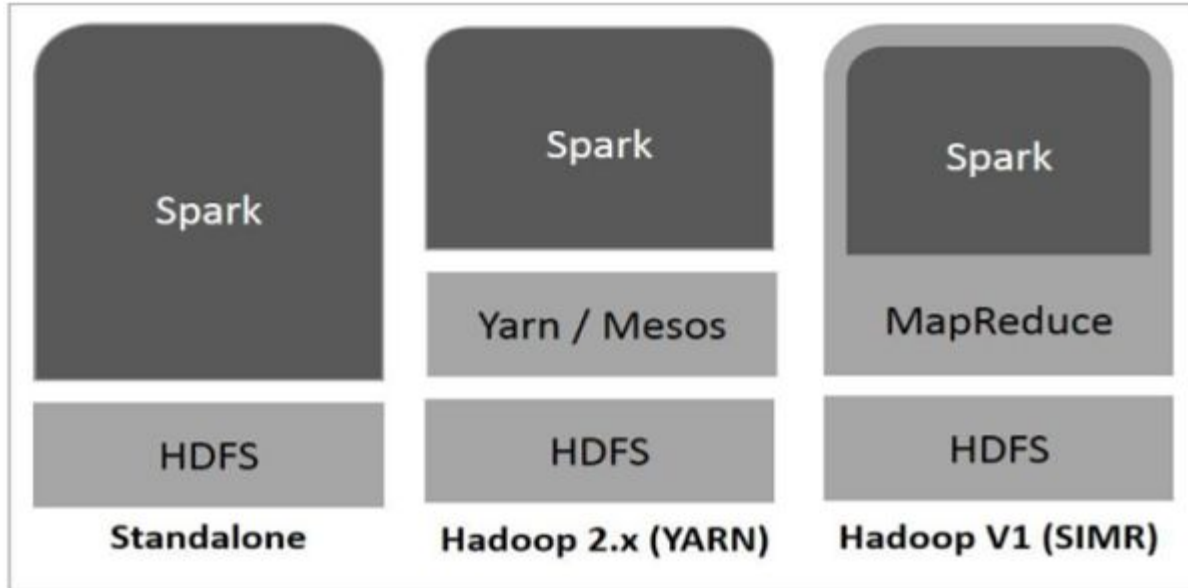
Spark

- Spark was introduced for speeding up the Hadoop computational computing software process.
- Resilient Distributed Datasets(RDD)
 - in-memory cluster computing

Spark




Spark



Spark

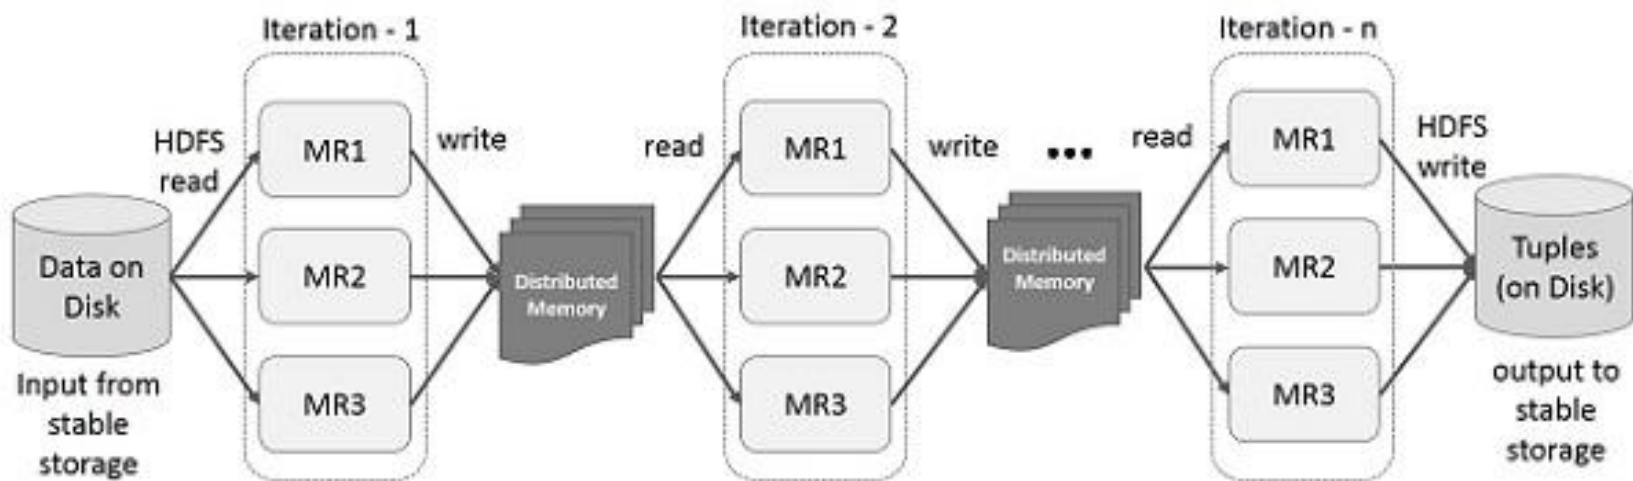
Hadoop mapreduce process



Most of the Hadoop applications, they spend more than 90% of the time doing HDFS read-write operations.

Spark

iterative operations on Spark RDD

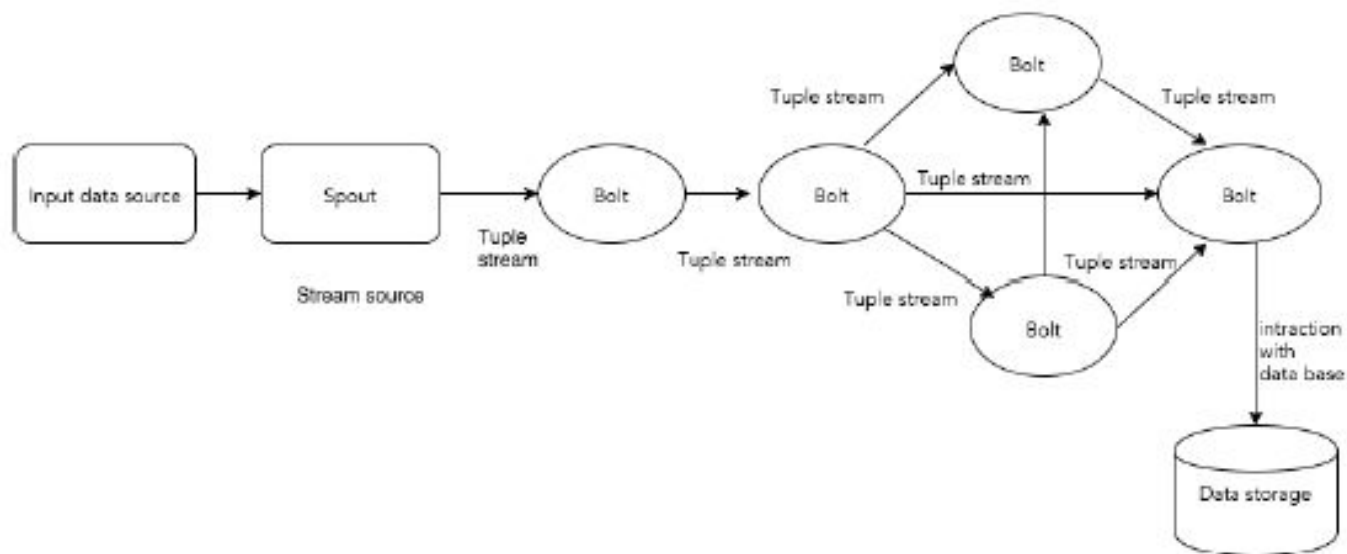


store intermediate results in a distributed memory instead of disk and make the system faster

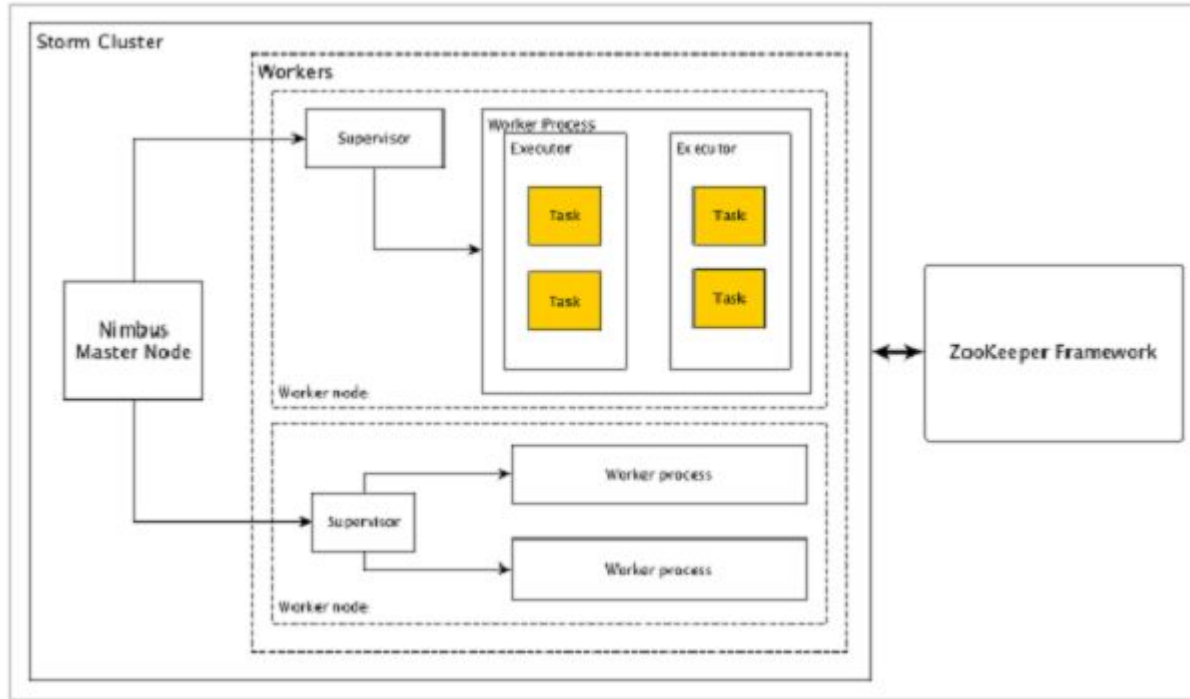
Storm

- Real-time stream processing
- A Storm streaming process can access tens of thousands messages per second on cluster.
- Local mode – This mode is used for development, testing, and debugging because it is the easiest way to see all the topology components working together. In this mode, we can adjust parameters that enable us to see how our topology runs in different Storm configuration environments. In Local mode, storm topologies run on the local machine in a single JVM.

Storm



Storm



Testing Environment

Local host:

1.3 GHz Intel Core i5

8 GB 1600 MHz DDR3

121.3 GB APPLE SSD SD0128F

Package:

Apache hadoop/3.1.0 in pseudo distributed mode

mahout/0.13.0

jdk/1.8.0_171

Apache spark/2.3.1 in Local mode

scala/2.12.6

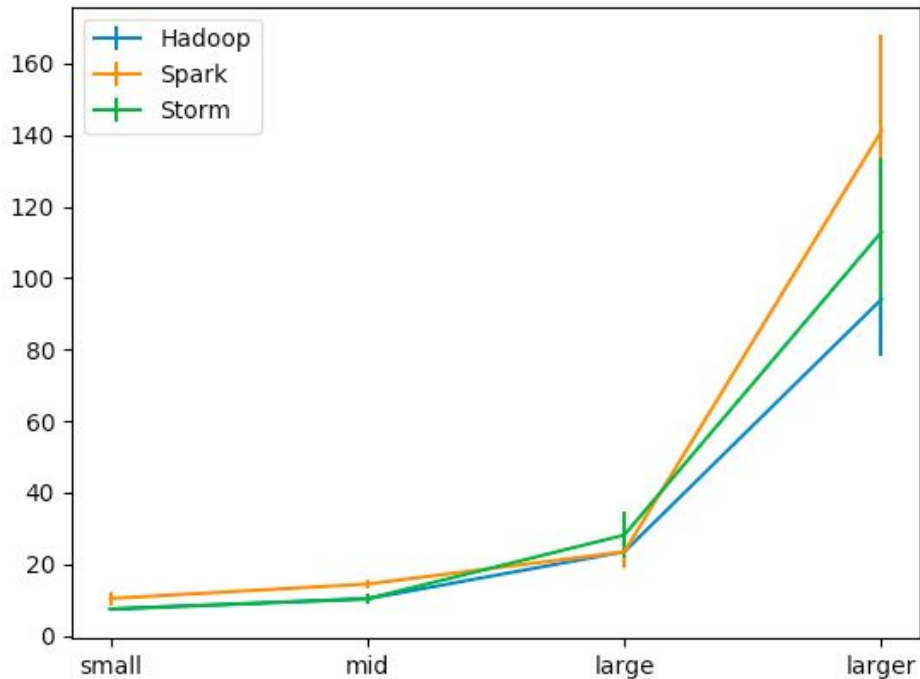
python/3.6.5

Apache storm/1.22 in standalone mode


zookeeper/3.4.10

Read/ Write

Write



Read



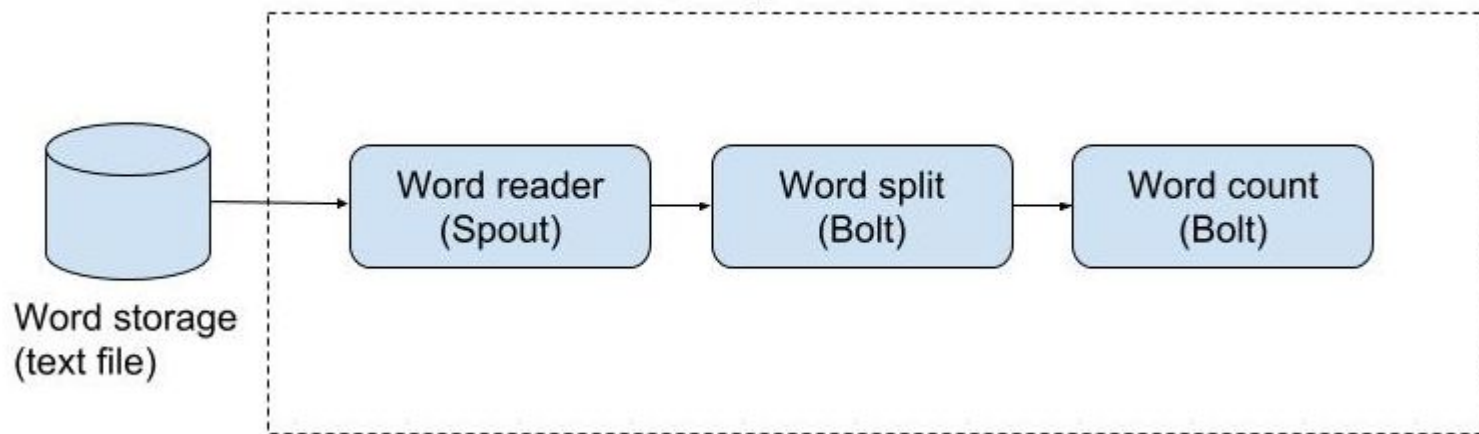
Time (sec) consumption on different size of data. Small :285KB, Medium :240MB, Large: 756MB, Larger: 5.33GB

Word Count

- Accumulate the frequency of each word
- Information retrieving like search engine
- Small Data: word_count.txt (58 KB)
Medium Data: t8_shakespeare (5.5 MB)
Large Data: dewiki_page_meta.xml (756 MB)
Larger Data: enwiktionary.xml (5.33 GB)

Storm

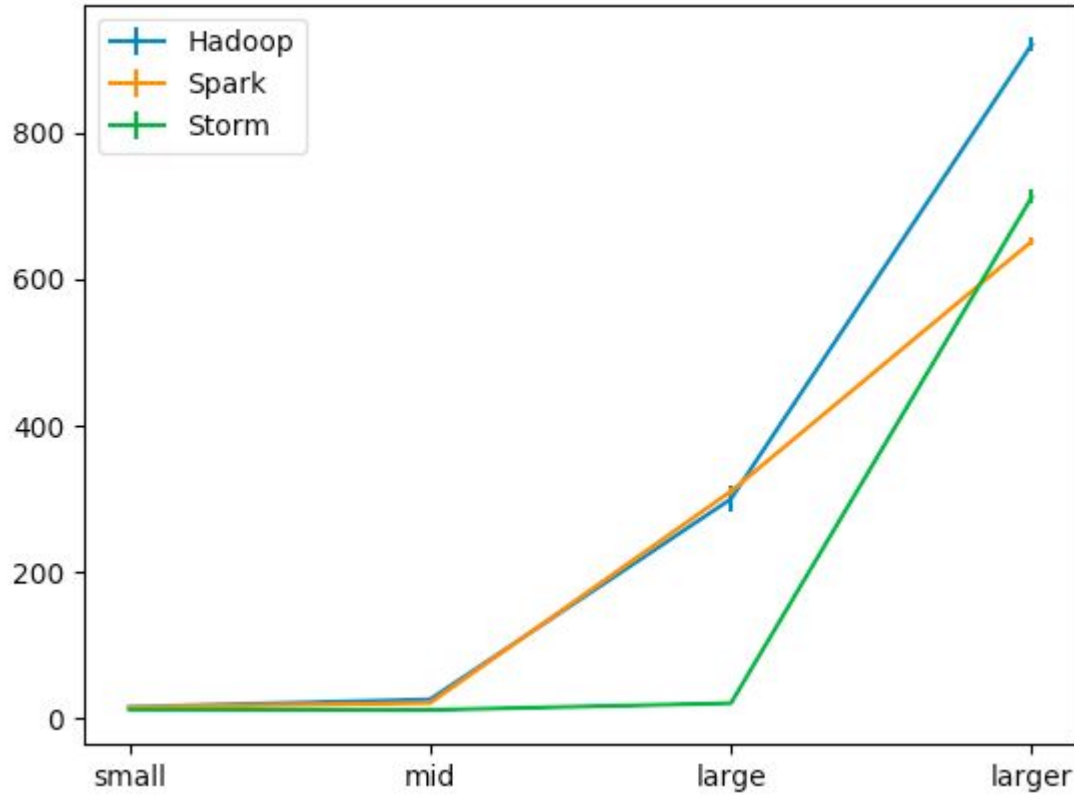
Topology



//Topology definition

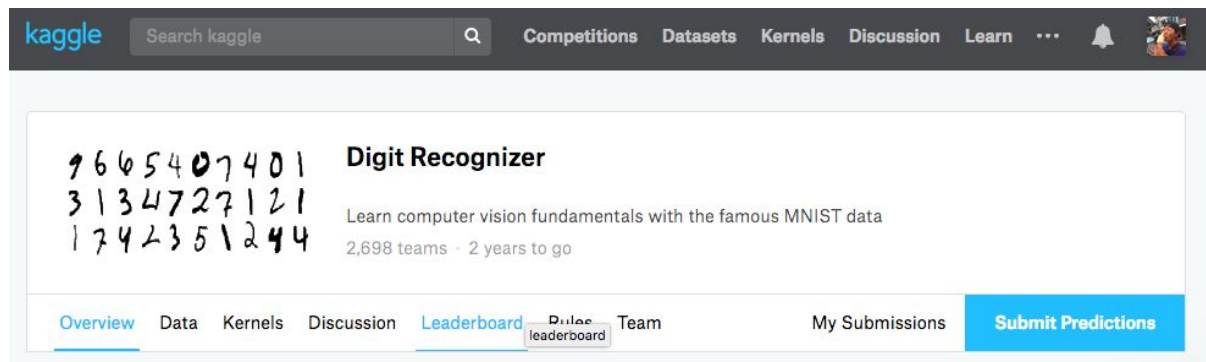
```
TopologyBuilder builder = new TopologyBuilder();
builder.setSpout("word-reader", new WordReader());
builder.setBolt("word-normalizer", new WordNormalizer())
    .shuffleGrouping("word-reader");
builder.setBolt("word-counter", new WordCounter(), 1)
    .fieldsGrouping("word-normalizer", new Fields("word"));
```

Word Count



Time consumption
(sec) on different size
of text data

K-means



The screenshot shows the Kaggle website header with the logo and navigation links: Search kaggle, Competitions, Datasets, Kernels, Discussion, Learn, and a user profile icon. Below the header is the competition banner for 'Digit Recognizer'. The banner features a 3x3 grid of handwritten digits: 9, 6, 6, 5, 4, 0, 7, 4, 0, 1 in the first row; 3, 1, 3, 4, 7, 2, 7, 1, 2, 1 in the second row; and 1, 7, 4, 2, 3, 5, 1, 2, 4, 4 in the third row. To the right of the digits, the text reads 'Digit Recognizer', 'Learn computer vision fundamentals with the famous MNIST data', and '2,698 teams · 2 years to go'. At the bottom of the banner is a navigation bar with links: Overview, Data, Kernels, Discussion, Leaderboard (highlighted), and a dropdown menu for 'Public leaderboard' and 'Team'. To the right of the navigation bar are links for 'My Submissions' and a blue 'Submit Predictions' button.

- <https://www.kaggle.com/c/digit-recognizer/data>

Given raw data of image and label, predict the unlabeled data


- No need label when using k-means

- Test_little.csv: 3600 data points x 783 pixels


Test.csv: 14000 data points

Train.csv: 25500 data points


K-means: test_little



K-means: test



K-means: train



Time consumption
(sec) on different k
value

Conclusion & Future work

- Easy development environment excepted some package dependency
- Most of our experiments run under single host mode
- Other performance evaluation of file system
 - Scalability: When # of clients increases
 - Fault tolerance: server & host
- Other applications: Neural network